

AEGIS: Mining Graph Structure for Adversarial Vulnerability Analysis of Graph Neural Networks

Anonymous Author(s)

Abstract—Graph neural networks are increasingly deployed in safety-critical domains — financial fraud detection, drug interaction prediction, infrastructure monitoring — where adversarial manipulation of graph structure can cause silent, hard-to-detect failures. Yet practitioners currently lack tools to answer a basic question before deployment: *which edges, if perturbed, would flip this model’s predictions?* We introduce AEGIS (Adversarial Evaluation of Graph Integrity via Sensitivity), a framework that answers this question by constructing a constrained sensitivity matrix S_c mapping edge perturbations to hidden-state shifts. From S_c , AEGIS extracts the mathematically optimal attack direction (via SVD), per-edge vulnerability rankings, and per-node first-order sensitivity radii — all in a single forward-backward pass. The computation applies to any differentiable GNN; for contractive implicit models (IGNN-class), it additionally provides a critical perturbation budget $\varepsilon_{\text{crit}}$ and convergence guarantees. Under realistic (symmetric, edge-only) perturbations, first-order predictions achieve tightness 1.00 ± 0.01 at $\varepsilon = 0.01$ and remain within 15% at $\varepsilon = 0.10$ across 7 architectures and 9 datasets. The SVD-optimal attack inflicts 2–8 $\times$ more damage than random perturbation. On IEEE power grids, the vulnerability spectrum recovers N-1 contingency rankings (P@10 = 0.66–0.81) without domain-specific inputs, demonstrating that structural sensitivity analysis can serve as a general-purpose pre-deployment diagnostic. Code will be released upon publication.

Index Terms—Graph neural networks, adversarial robustness, structural sensitivity, deep equilibrium models, graph perturbation, implicit function theorem

I. INTRODUCTION

Graph neural networks are increasingly deployed in domains where incorrect predictions carry tangible consequences: fraudulent accounts evade detection when an adversary inserts edges into a financial transaction graph [1], drug interaction models produce unsafe recommendations when molecular graphs are perturbed [2], and power grid monitoring systems may miss critical contingencies if their graph-based representations are structurally fragile [3]. In all these settings, the graph topology is not merely an input format but a potential attack surface. A practitioner preparing to deploy a GNN faces a fundamental question that current tools leave unanswered: *which edges in this graph, if perturbed, would cause the model’s predictions to fail — and by how much?*

This question is distinct from both adversarial attack and certified defense. Attack methods such as Netattack [1] and Mettack [4] search for damaging perturbations using surrogate gradients or meta-learning, but they provide no guarantee that the perturbation found is optimal, and they do not quantify per-node vulnerability. Certified defenses based on randomized smoothing [5], [6] or interval bound propagation [7],

[8], [9] provide robustness guarantees, but they are uniform: every node receives the same certificate, revealing nothing about *which* edges or nodes are structurally weak. Empirical defenses [10], [11], [12], [13] harden the model but offer no formal understanding of residual vulnerability. What is missing is a *structural vulnerability map* — a tool that, given a trained GNN and a graph, identifies the precise attack surface: the mathematically optimal perturbation direction, the weakest edges, and the perturbation tolerance of each node.

We introduce AEGIS (Adversarial Evaluation of Graph Integrity via Sensitivity), a framework that produces exactly this map. The central object is the *constrained sensitivity matrix* $S_c \in \mathbb{R}^{D \times |E|}$, which projects the full N^2 -dimensional adjacency perturbation space onto the $|E|$ -dimensional space of realistic (symmetric, edge-only) perturbations. From S_c , three outputs follow directly: (i) the SVD yields the first-order optimal attack direction, (ii) column norms give per-edge vulnerability scores, and (iii) the classification margin divided by the sensitivity norm gives per-node tolerance radii. The S_c computation applies to any differentiable GNN via Jacobian evaluation (theorem 4). For contractive implicit GNNs (IGNN-class [14]), the implicit function theorem additionally provides a critical perturbation budget $\varepsilon_{\text{crit}} = (1 - \kappa) / \|W\|_2$ and a three-regime characterization of model behavior under increasing perturbation (theorem 1).

The practical value is that these predictions are quantitatively accurate: constrained first-order tightness is 1.00 ± 0.01 at $\varepsilon = 0.01$ and remains within 15% at $\varepsilon = 0.10$, across 7 GNN architectures (IGNN, GCN, GIN, GAT, GraphSAGE, APPNP) and 9 datasets spanning citation networks, e-commerce, encyclopedias, and power grids. On IEEE power grids, the vulnerability spectrum independently recovers N-1 contingency rankings (P@10 = 0.66–0.81), demonstrating that structural sensitivity analysis transfers to real-world engineering problems without domain-specific modifications.

Contributions.

- 1) **Constrained sensitivity framework.** We derive the constrained sensitivity matrix S_c for any differentiable GNN (theorem 4). For contractive implicit models, we additionally prove a critical budget $\varepsilon_{\text{crit}}$ and three vulnerability regimes (theorem 1).
- 2) **Unified vulnerability analysis across architectures.** From S_c we extract SVD-optimal attacks (2–8 $\times$ stronger than random), per-edge vulnerability rankings, and per-node first-order sensitivity radii, validated on 7 architectures (section V-H).

- 3) **Cross-domain validation and case study.** We validate across 9 datasets in 4 domains, including a power grid case study where AEGIS recovers N-1 contingency rankings and reveals implicit physics consistency of the equilibrium architecture (section VI).

II. PRELIMINARIES AND THREAT MODEL

This section establishes the notation, formalizes the model class under consideration, specifies the threat model, and states the problem that AEGIS addresses.

A. Notation and Formalization

Let $G = (V, E, A)$ denote an undirected graph with $N = |V|$ nodes, edge set E , and adjacency matrix $A \in \mathbb{R}^{N \times N}$. Each node $v \in V$ carries a feature vector $x_v \in \mathbb{R}^{d_{\text{in}}}$, collected into $X \in \mathbb{R}^{N \times d_{\text{in}}}$. We write $\hat{A} = D^{-1/2}(A + I)D^{-1/2}$ for the symmetric normalized adjacency with self-loops, where D is the degree matrix of $A + I$.

Deep equilibrium GNNs. A DEQ-GNN [15], [16] defines a weight-tied operator $F_\theta : \mathbb{R}^{N \times d} \times \mathbb{R}^{N \times N} \rightarrow \mathbb{R}^{N \times d}$ and iterates it from $Z_0 = 0$ until convergence:

$$Z_{k+1} = F_\theta(Z_k, A), \quad z^* = \lim_{k \rightarrow \infty} Z_k. \quad (1)$$

When F_θ is contractive in operator norm ($\kappa = \|J_z\|_2 < 1$, where $J_z = \partial F / \partial z|_{z^*}$), the Banach fixed-point theorem guarantees a unique equilibrium $z^* = F_\theta(z^*, A)$ regardless of initialization.

IGNN instantiation. The Implicit Graph Neural Network (IGNN) [14] uses the operator

$$F(Z, A) = \sigma(\hat{A}ZW^\top + X_{\text{proj}}), \quad (2)$$

where $W \in \mathbb{R}^{d \times d}$ is the shared weight matrix, $X_{\text{proj}} \in \mathbb{R}^{N \times d}$ is a learned feature projection, and σ is ReLU activation. Contractivity is enforced by constraining $\|W\|_2$ via spectral normalization [17], ensuring $\kappa \leq \|\hat{A}\|_2 \cdot \|W\|_2 < 1$. Training uses implicit differentiation through the fixed point [18], [19], and Jacobian regularization [20] can further stabilize convergence. A linear classification head $f(z^*) = W_{\text{cls}}z^* + b$ maps the equilibrium to class predictions.

Implicit function theorem. At the equilibrium, the equation $G(z^*, A) = z^* - F(z^*, A) = 0$ defines z^* as an implicit function of A . The IFT gives the sensitivity of z^* to any parameter p :

$$\frac{\partial z^*}{\partial p} = (I - J_z)^{-1} \frac{\partial F}{\partial p} \Big|_{z^*}, \quad (3)$$

where the resolvent $(I - J_z)^{-1}$ exists and satisfies $\|(I - J_z)^{-1}\|_2 \leq 1/(1 - \kappa)$ by the Neumann series (convergent because $\kappa < 1$). This bound uses the operator norm $\kappa = \|J_z\|_2$; the weaker spectral-radius bound $1/(1 - \rho)$ holds only for normal J_z . The pseudospectral index $\eta = \|(I - J_z)^{-1}\|_2 \cdot (1 - \rho) \geq 1$ [21] quantifies this gap, with $\eta = 1$ for normal matrices. The IFT has been used for hyperparameter optimization [22] and implicit differentiation [18], but not for adversarial structural analysis.

B. Threat Model

We consider a white-box attacker who has full access to the trained IGNN model and perturbs the normalized adjacency matrix:

$$\hat{A}' = \hat{A} + \delta \hat{A}, \quad \text{subject to} \quad \|\delta \hat{A}\|_F \leq \varepsilon. \quad (4)$$

The perturbation is further constrained to be:

- **Symmetric:** $\delta \hat{A} = \delta \hat{A}^\top$ (preserving the undirected nature of the graph).
- **Edge-only:** $[\delta \hat{A}]_{ij} = 0$ for $(i, j) \notin E$ (only existing edge weights may be modified; no new edges are introduced).
- **Continuous:** edge weights are perturbed continuously in $\mathbb{R}$, not discretely flipped.

This threat model captures continuous edge-weight perturbations to the normalized message-passing operator. It aligns naturally with the IGNN architecture, where $\hat{A}$ appears directly in F (eq. (2)). Discrete edge insertions or deletions, which change the graph topology and require recomputing the degree normalization $D^{-1/2}$, are outside the formal guarantee.

The key distinction from input-feature perturbation (studied for DEQs in [23], [24]) is that structural perturbation changes the *operator itself*: the equilibrium depends on A through both the state Jacobian $J_z = \text{diag}(\sigma') \cdot (\hat{A} \otimes W)$ and the adjacency Jacobian $J_A = \partial F / \partial \text{vec}(A)$. This dual dependence makes structural sensitivity qualitatively different from input sensitivity.

C. Problem Statement

Given a trained contractive implicit GNN with equilibrium z^* on graph G , we aim to answer three questions:

- 1) **Vulnerability ranking:** Which edges, if perturbed, cause the largest shift in the equilibrium? That is, find a ranking of edges $(i, j) \in E$ by their adversarial impact on z^* .
- 2) **Optimal attack:** What is the worst-case perturbation $\delta \hat{A}^*$ that maximizes $\|\Delta z^*\|_F$ subject to the threat model constraints in eq. (4)?
- 3) **Robustness assessment:** For each node v , what is the largest perturbation budget r_v such that any $\|\delta \hat{A}\|_F < r_v$ preserves the classification of v ?

AEGIS addresses all three questions through a unified object: the constrained sensitivity matrix S_c derived from the IFT applied to the equilibrium equation. The following sections formalize this construction (section III) and describe the computational pipeline (section IV).

III. ADVERSARIAL EQUILIBRIUM THEORY

Having formalized the problem in section II, we now derive the theoretical foundation for answering the three questions posed there. The main result (theorem 1) characterizes vulnerability regimes as a function of perturbation magnitude, while the constrained sensitivity matrix S_c provides the unified object from which vulnerability rankings, optimal attacks, and robustness radii are all extracted.

Theorem 1 (Vulnerability Characterization for Contractive Implicit GNNs). *Let $F_\theta(Z, A) = \sigma(AZW^\top + X_{\text{proj}})$ be an IGNN-class operator satisfying:*

- (A1) σ is ReLU (or any 1-Lipschitz activation with $\sup |\sigma'| = 1$);
- (A2) W is spectral-norm constrained: $\|W\|_2 \leq c$;
- (A3) The state Jacobian satisfies $\|J_z\|_2 \leq \kappa < 1$ at the fixed point z^* (contraction in operator norm).

Define the critical budget

$$\varepsilon_{\text{crit}} = \frac{1 - \kappa}{\|W\|_2}. \quad (5)$$

Then structural perturbations δA with $\|\delta A\|_F = \varepsilon$ exhibit three regimes:

(a) **Subcritical** ($\varepsilon < \varepsilon_{\text{crit}}$): The perturbed operator remains contractive with a unique fixed point satisfying

$$\|\Delta z^*\|_F \leq \sigma_1(S) \cdot \varepsilon + O(\varepsilon^2), \quad (6)$$

where $S = (I - J_z)^{-1} J_A$ is the structural sensitivity matrix and $\sigma_1(S) \leq \|J_A\|_{\text{op}} / (1 - \kappa)$.

(b) **Critical** ($\varepsilon \rightarrow \varepsilon_{\text{crit}}$): The resolvent norm $\|(I - J'_z)^{-1}\|_2$ diverges as $\Omega(1/(\varepsilon_{\text{crit}} - \varepsilon))$. Intuitively, as the perturbation approaches the critical budget, the system's sensitivity to further perturbation grows without bound.

(c) **Supercritical** ($\varepsilon > \varepsilon_{\text{crit}}$): The sufficient contraction certificate becomes void ($\|J'_z\|_2$ may exceed 1). The perturbed operator may still be contractive depending on perturbation direction, but the Banach fixed-point theorem no longer guarantees uniqueness, and the first-order guarantees from part (a) cease to apply.

The intuition behind this three-regime picture is physical: a contractive system has a restoring force that pulls trajectories toward the equilibrium. Small perturbations shift the equilibrium smoothly (subcritical). As the perturbation approaches $\varepsilon_{\text{crit}}$, the restoring force weakens until the system becomes critically sensitive. Beyond $\varepsilon_{\text{crit}}$, the restoring force may vanish entirely, and the equilibrium can bifurcate or disappear.

Proof. **Part (a).** The Jacobian of F with respect to $\text{vec}(Z)$ takes the form $J_z = \text{diag}(\sigma') \cdot (\hat{A} \otimes W)$, where $\sigma' \in \{0, 1\}^{Nd}$ for ReLU (A1). The operator norm satisfies $\kappa = \|J_z\|_2 \leq \|\hat{A}\|_2 \cdot \|W\|_2 \cdot \sup |\sigma'| = \|\hat{A}\|_2 \cdot \|W\|_2$. Applying the IFT to $G(z, A) = z - F(z, A) = 0$ at (z^*, A) :

$$\Delta z^* = (I - J_z)^{-1} J_A \cdot \text{vec}(\delta A) + O(\|\delta A\|^2). \quad (7)$$

Taking operator norms and using the Neumann series bound $\|(I - J_z)^{-1}\|_2 \leq 1/(1 - \kappa)$ (convergent because $\kappa < 1$) yields (6).

For contractivity preservation: the perturbed Jacobian satisfies $\|J'_z\|_2 \leq (\|A\|_2 + \|\delta A\|_F) \cdot \|W\|_2$ (using $\|\cdot\|_2 \leq \|\cdot\|_F$ for the perturbation). The condition $\varepsilon < (1 - \kappa)/\|W\|_2$ ensures $\|J'_z\|_2 < 1$.

Part (b). As $\varepsilon \rightarrow \varepsilon_{\text{crit}}$, we have $\|J'_z\|_2 \rightarrow 1$, so $\|(I - J'_z)^{-1}\|_2 \geq 1/(1 - \|J'_z\|_2) \rightarrow \infty$. Unlike the spectral-radius bound $1/(1 - \rho)$, which requires normality, the operator-norm bound holds for all J_z with $\kappa < 1$. In our experiments,

the pseudospectral index η ranges from 1.02 to 1.28 (section V-F), confirming mild non-normality.

Part (c). When $\|J'_z\|_2 \geq 1$, the contraction mapping condition fails. The Banach theorem no longer guarantees existence or uniqueness; the iteration may converge to a different equilibrium, oscillate, or diverge. $\square$

Remark (conservativeness). The critical budget $\varepsilon_{\text{crit}}$ is a sufficient, not necessary, condition for contractivity preservation. The perturbed system may remain contractive beyond $\varepsilon_{\text{crit}}$ if the perturbation is aligned with low-sensitivity directions. The unconstrained bound also uses worst-case $\|\cdot\|_F$ relaxation. Under constrained perturbations, the first-order prediction achieves tightness ≈ 1.00 (section V-A). We use $\kappa = \|J_z\|_2$ rather than the spectral radius $\rho(J_z)$ because the Neumann series requires operator-norm contractivity [25]; for normal J_z the two coincide, while for non-normal J_z the κ -based bound is more conservative but requires no normality assumption.

Proposition 2 (Optimal First-Order Structural Attack). *The perturbation δA^* maximizing $\|\Delta z^*\|$ to first order subject to $\|\delta A\|_F \leq \varepsilon$ is $\delta A^* = \varepsilon \cdot \text{reshape}(v_1, N \times N)$, where v_1 is the leading right singular vector of S . The maximum first-order shift is $\varepsilon \cdot \sigma_1(S)$.*

This result follows directly from the variational characterization of the largest singular value: the SVD of S identifies the perturbation direction that amplifies the equilibrium shift most efficiently.

Threat model and constrained sensitivity. As discussed in section II, realistic graph perturbations are symmetric and edge-only. We perturb the normalized adjacency $\hat{A}$ directly, subject to $\|\delta \hat{A}\|_F \leq \varepsilon$, symmetry ($\delta \hat{A} = \delta \hat{A}^\top$), and support restricted to existing edges. Discrete edge additions that change topology are outside the formal guarantee.

To enforce these constraints, we construct the *constrained sensitivity matrix* $S_c \in \mathbb{R}^{D \times |E|}$ with column

$$[S_c]_{:,k} = S_{:,iN+j} + S_{:,jN+i} \quad (8)$$

for each edge $(i, j) \in E$. This reduces the perturbation space from N^2 dimensions to $|E|$ dimensions and enforces symmetry by construction. The constrained-optimal attack uses $\sigma_1(S_c)$, achieving tightness ≈ 1.00 empirically (section V-A). The construction of S_c is the central technical contribution of this work: it transforms unconstrained sensitivity analysis into a tight, practical tool for realistic graph perturbations.

Proposition 3 (Per-Node First-Order Sensitivity Radius). *Assume a linear classification head $f(z) = Wz + b$ (as in standard IGNN). For node v with margin $m_v = f_{y_v}(z_v^*) - \max_{c \neq y_v} f_c(z_v^*) > 0$, the first-order sensitivity radius is*

$$r_v = \frac{m_v}{\|W_{y_v} - W_{c^*}\|_2 \cdot \|S_v\|_2}, \quad (9)$$

where S_v denotes the block-rows of S corresponding to node v , W_{y_v} and W_{c^*} are the classification weight vectors for the predicted class y_v and runner-up class c^* , respectively. Any

perturbation with $\|\delta\hat{A}\|_F < r_v$ preserves the classification of node v to first order.

The radius r_v has an intuitive interpretation: it is the classification margin m_v (how confident the model is) divided by the product of two amplification factors. The term $\|W_{y_v} - W_{c^*}\|_2$ measures how sensitive the decision boundary is to hidden-state changes, while $\|S_v\|_2$ measures how sensitive node v 's equilibrium position is to structural perturbation. Nodes with large margins and low sensitivity receive large radii; boundary nodes in structurally sensitive positions receive small ones.

These sensitivity radii are deterministic (not probabilistic), unlike randomized smoothing [5] which holds with probability $1 - \alpha$. They are first-order quantities: locally tight for small ε but increasingly conservative as perturbation magnitude grows. Without bounding the second-order remainder $O(\|\delta A\|_F^2)$, they are local sensitivity measures rather than global robustness certificates. section V-C verifies empirically that the radii hold (0% breach rate at $\varepsilon = 0.01$, $< 1\%$ at $\varepsilon = 0.10$).

A. Generalization to Explicit GNNs

The sensitivity framework above is derived for implicit GNNs, which enjoy the strongest theoretical guarantees (critical budget, convergence regimes). The S_c construction itself, however, generalizes to any differentiable K -layer GNN as a computational tool for vulnerability analysis.

Proposition 4 (Structural Sensitivity for K -Layer GNNs). *Let $Z_K = g_K \circ \dots \circ g_1(X, A)$ be a K -layer GNN where each layer g_l is differentiable with respect to both its input and A . Define the unrolled sensitivity matrix*

$$S_K = \frac{\partial \text{vec}(Z_K)}{\partial \text{vec}(A)} = \sum_{l=1}^K \left(\prod_{k=l+1}^K J_z^{(k)} \right) J_A^{(l)}, \quad (10)$$

where $J_z^{(k)} = \partial g_k / \partial Z_{k-1}$ and $J_A^{(l)} = \partial g_l / \partial \text{vec}(A)$. Then:

(a) *The first-order shift bound $\|\Delta Z_K\|_F \leq \sigma_1(S_K) \cdot \|\delta A\|_F + O(\|\delta A\|^2)$ holds, with*

$$\sigma_1(S_K) \leq \sum_{l=1}^K \left(\prod_{k=l+1}^K \|J_z^{(k)}\|_2 \right) \|J_A^{(l)}\|_2. \quad (11)$$

(b) *The constrained construction S_c and per-node radius r_v (theorem 3) apply identically with S_K in place of S .*

For weight-tied models ($J_z^{(k)} = J_z$, $J_A^{(l)} = J_A$ for all k, l), the bound simplifies to $\sigma_1(S_K) \leq \|J_A\|_2 \cdot (1 - \kappa^K)/(1 - \kappa)$, which converges to the IGNN bound $\|J_A\|_2/(1 - \kappa)$ as $K \rightarrow \infty$.

What explicit GNNs gain and lack. From theorem 4, any differentiable GNN inherits: (i) the S_c construction, (ii) SVD-optimal attacks, (iii) per-edge vulnerability rankings, and (iv) per-node first-order sensitivity radii. What explicit models *lack* is the critical budget $\varepsilon_{\text{crit}}$ and the three-regime convergence characterization of theorem 1, which require the contraction assumption (A3). section V-H validates this generalization on six explicit architectures.




Fig. 1: The AEGIS pipeline. Data flows top to bottom; annotations show intermediate quantities passed between stages. The bottleneck is Stage 2 ($O(D^2)$ backward passes).

IV. AEGIS CONSTRUCTION

The theoretical results of section III provide the mathematical foundation for vulnerability analysis, but translating them into a practical tool requires careful engineering. This section describes the AEGIS pipeline: how its components fit together, what each one computes, and how the system scales to real graphs.

A. Architecture Overview

AEGIS takes as input a trained GNN (any differentiable architecture) together with a target graph and produces three outputs: (i) a per-edge vulnerability spectrum ranking edges by adversarial impact, (ii) SVD-optimal attack perturbations, and (iii) per-node first-order sensitivity radii. For contractive implicit GNNs (IGNN-class satisfying A1–A3), the pipeline additionally computes the critical budget $\varepsilon_{\text{crit}}$ and convergence diagnostics; these theoretical guarantees are specific to the implicit case. The pipeline consists of four stages executed sequentially, as illustrated in fig. 1.

B. Stage 1: Subgraph Extraction and Forward Pass

Full-graph analysis is infeasible for large graphs because the sensitivity matrix has dimension $D \times N^2$ where $D = N \cdot d$. AEGIS addresses this by extracting a local BFS ego-subgraph of N nodes centered on the target node. The model is then run on this subgraph: for implicit GNNs, this means iterating to the fixed point z^* ; for explicit K -layer GNNs, this is a

standard forward pass yielding Z_K . Localization is justified by the locality of per-node vulnerability: a node’s robustness radius depends on its own margin and the sensitivity of its immediate neighborhood (section V-E).

C. Stage 2: Sensitivity Computation

AEGIS computes the structural sensitivity matrix $S = \partial \text{vec}(Z)/\partial \text{vec}(A) \in \mathbb{R}^{D \times N^2}$, which maps adjacency perturbations to hidden-state shifts. The computation differs by architecture:

Implicit GNNs (IGNN-class): $S = (I - J_z)^{-1} J_A$, where $J_z = \partial F/\partial Z$ is the state Jacobian and $J_A = \partial F/\partial \text{vec}(A)$ is the adjacency Jacobian. Both are computed via row-by-row backward passes, followed by a linear solve. For IGNN with ReLU, $J_z = \text{diag}(\sigma') \cdot (\hat{A} \otimes W)$.

Explicit GNNs (K-layer): $S_K = \partial \text{vec}(Z_K)/\partial \text{vec}(A)$ is computed directly via finite differences on the forward pass. Each column of S_K measures the effect of perturbing one adjacency entry.

Both paths use automatic differentiation [26] with cost $O(D \cdot N^2)$ for the full sensitivity matrix.

D. Stage 3: Constrained Projection

Given the full sensitivity matrix $S \in \mathbb{R}^{D \times N^2}$ (from Stage 2, via either the linear solve for implicit GNNs or direct computation for explicit GNNs), AEGIS projects it into the *constrained sensitivity matrix* $S_c \in \mathbb{R}^{D \times |E|}$. Realistic graph perturbations are symmetric (undirected graphs) and supported only on existing edges. The k -th column of S_c (for edge $(i, j) \in E$) is

$$[S_c]_{:,k} = S_{:,iN+j} + S_{:,jN+i}. \quad (12)$$

The summation enforces symmetry: perturbing edge (i, j) affects both A_{ij} and A_{ji} simultaneously. This projection reduces the perturbation space from N^2 dimensions to $|E|$ dimensions and is what enables tight first-order predictions across all architectures (sections V-A and V-H).

E. Stage 4: Vulnerability Analysis

From S_c , AEGIS extracts three complementary outputs:

Optimal attack direction. The SVD of $S_c = U\Sigma V^\top$ yields the first-order optimal perturbation as $\delta A^* = \varepsilon \cdot v_1$ (the leading right singular vector of S_c , reshaped to edge weights). This perturbation maximizes the fixed-point shift $\|\Delta z^*\|_F$ subject to the Frobenius-norm budget ε , and achieves a shift of $\varepsilon \cdot \sigma_1(S_c)$.

Per-edge vulnerability spectrum. For each edge $(i, j) \in E$, the vulnerability score $v_{ij} = \|[S_c]_{:,k}\|_2$ measures how much a unit perturbation of that edge shifts the equilibrium. Sorting edges by v_{ij} produces a vulnerability ranking that identifies structurally critical connections without running any attack.

Per-node first-order sensitivity radii. For node v with classification margin m_v (gap between the predicted and

runner-up class logits), the sensitivity radius from theorem 3 is

$$r_v = \frac{m_v}{\|W_{y_v} - W_{c^*}\|_2 \cdot \|S_v\|_2}, \quad (13)$$

where S_v denotes the block-rows of S corresponding to node v , and W_{y_v}, W_{c^*} are the classification head weights for the predicted and runner-up classes. Any perturbation with $\|\delta \hat{A}\|_F < r_v$ preserves the classification of node v to first order (this is a local sensitivity guarantee, not a global certificate).

F. Computational Cost

The dominant cost is Stage 2 (Jacobian computation) at $O(D \cdot N^2)$ for J_A , followed by Stage 3 (linear solve) at $O(D^3)$. For the default subgraph size $N = 50$ with $d = 64$ (giving $D = 3,200$), the full pipeline completes in 1.3 seconds on a single GPU. The practical limit is $N \approx 200$ ($D = 12,800$, 8.1 seconds), requiring approximately 6.5 GB GPU memory for J_z (D^2 entries) and J_A ($D \times N^2$ entries) in float32; see section V-D for detailed timings.

V. EXPERIMENTS

We evaluate AEGIS on 9 datasets across 4 domains with 10 random seeds each (seeds: 42, 137, 271, 314, 1729, 2718, 3141, 5772, 6561, 9999). All experiments are implemented in PyTorch [26]. Graph benchmark experiments use IGNN [14] with hidden dimension $d = 64$, spectral-norm constrained weights [17], and early stopping on validation accuracy over 200 epochs. Power flow experiments use ContractiveGCN-PF, an IGNN variant for AC power flow (section VI).

Datasets. We use five graph benchmarks: Cora (2,708 nodes), Citeseer (3,327), and Pubmed (19,717) from the Planetoid collection [27]; Amazon Photo (7,650) from the Amazon co-purchase graphs [28]; and WikiCS (11,701) from Wikipedia-based benchmarks [29]. For power flow, we use four IEEE standard test cases [30]: case14 (14 buses), case30 (30), case57 (57), and case118 (118), with training data generated via PandaPower [31].

Notation. All formal bounds (theorem 1) use the operator-norm contraction constant $\kappa = \|J_z\|_2$. Tables report the spectral radius $\rho(J_z)$ as a diagnostic quantity. Since non-normality is mild ($\eta = 1.02$ – 1.28 , section V-F), we have $\kappa \approx \rho$ in practice; the $\varepsilon_{\text{crit}}$ values computed with ρ are at most 28% optimistic relative to the κ -based formal threshold.

A. Cross-Domain Adversarial Analysis

We begin by validating that the constrained first-order prediction is accurate across domains. table I reports the tightness ratio (actual equilibrium shift divided by predicted shift at $\varepsilon = 0.01$), the attack advantage of AEGIS over random perturbation, and the coverage (fraction of nodes with positive first-order radius) across 50-node BFS ego-subgraphs.

Tightness is 1.00 ± 0.01 across all datasets at $\varepsilon = 0.01$. First-order accuracy at small ε is mathematically expected; the contribution is that S_c makes this tractable under *constrained* perturbations ($N^2 \rightarrow |E|$ projection). table II shows how

TABLE I: Cross-domain adversarial analysis (IGNN, 10 seeds). ρ : spectral radius of state Jacobian. Tight.: actual/predicted shift ratio. AtkAdv: AEGIS damage / random damage. $\varepsilon_{\text{crit}}$: critical perturbation budget (theorem 1). Test Acc / Cert%: full-graph / subgraph accuracy (%).

Dataset	Test Acc	$\rho(J_z)$	Tight.	AtkAdv	$\varepsilon_{\text{crit}}$	Cert%
Cora	77.5 $\pm$ 1.7	.33 $\pm$.14	1.01 $\pm$.01	3.6 $\pm$.6	.66 $\pm$.15	81 $\pm$ 9
Citeseer	66.0 $\pm$ 0.7	.59 $\pm$.02	1.01 $\pm$.01	4.1 $\pm$.5	.41 $\pm$.02	83 $\pm$ 4
Pubmed	78.9 $\pm$ 0.7	.41 $\pm$.11	1.01 $\pm$.01	3.2 $\pm$.4	.59 $\pm$.11	74 $\pm$ 23
Amazon	94.8 $\pm$ 0.3	.14 $\pm$.02	1.00 $\pm$.00	3.8 $\pm$.2	.86 $\pm$.02	92 $\pm$ 4
WikiCS	77.9 $\pm$ 0.4	.34 $\pm$.04	1.00 $\pm$.01	3.8 $\pm$.4	.66 $\pm$.04	70 $\pm$ 3

TABLE II: Tightness degradation with increasing ε (10 seeds, S_c -optimal constrained perturbation). Flip%: fraction of nodes whose prediction changes.

Dataset	$\varepsilon=.01$	$\varepsilon=.05$	$\varepsilon=.10$	$\varepsilon=.20$
<i>Tightness (actual/predicted shift)</i>				
Cora	1.01 $\pm$.00	1.07 $\pm$.01	1.15 $\pm$.03	1.36 $\pm$.08
Citeseer	1.01 $\pm$.00	1.07 $\pm$.01	1.16 $\pm$.03	1.39 $\pm$.07
WikiCS	1.00 $\pm$.01	1.01 $\pm$.01	1.02 $\pm$.01	1.05 $\pm$.01
<i>Prediction flip rate (%)</i>				
Cora	0.4 $\pm$ 0.8	0.6 $\pm$ 0.9	1.8 $\pm$ 2.4	5.8 $\pm$ 5.0
Citeseer	0.0 $\pm$ 0.0	0.6 $\pm$ 1.3	2.0 $\pm$ 2.4	3.0 $\pm$ 2.6
WikiCS	0.2 $\pm$ 0.6	0.4 $\pm$ 0.8	0.6 $\pm$ 0.9	1.2 $\pm$ 1.0

tightness degrades at larger budgets: the first-order prediction remains within 7% at $\varepsilon = 0.05$ and within 16% at $\varepsilon = 0.10$ on Cora/Citeseer, and within 5% at $\varepsilon = 0.20$ on WikiCS. Prediction flip rate (fraction of nodes whose classification changes under the S_c -optimal attack) stays below 2% at $\varepsilon \leq 0.10$ and below 6% at $\varepsilon = 0.20$. The attack advantage ($3.2\text{--}4.1\times$ over random in table I) confirms that S_c identifies the most vulnerable perturbation directions. Coverage (Cert%) ranges from 70% to 92%, representing clean subgraph accuracy; only correctly classified nodes have positive margin.

Surrogate baseline. We also compared against Mettack [4] (GCN [32] surrogate): AEGIS inflicts 3–10 $\times$ more damage at every budget (30/30 wins across $k = 1, \dots, 5$ edge removals). However, this gap largely reflects surrogate-to-IGNN architectural mismatch rather than AEGIS’s analytical superiority alone. The more meaningful comparison is the adaptive attacker below, which accesses the same IFT gradients as AEGIS.

B. Radius Comparison: AEGIS vs. Smoothing

We compare AEGIS sensitivity radii against randomized smoothing [5] ($\sigma \in \{0.01, 0.05, 0.10, 0.15\}$, 1000 MC samples, 10 seeds). At $\sigma \leq 0.10$, smoothing radii are *constant* across all nodes ($r = \sigma \cdot \Phi^{-1}(p_{\text{lower}})$) because all samples predict correctly, yielding larger absolute radii (0.123 vs. AEGIS’s 0.046 at $\sigma = 0.05$) but zero per-node differentiation. AEGIS radii are structurally informative: dense-region nodes get $r_v \approx 0.09$, boundary nodes $r_v \approx 0.01$. The two

TABLE III: Adaptive attack (white-box PGD via IFT gradients on IGNN) vs. AEGIS (10 seeds). Breach: fraction of certified nodes whose prediction changes. AEGIS/Adapt. dmg: mean equilibrium shift (ℓ_2). Ratio: Adapt./AEGIS damage.

Dataset	ε	Breach	AEGIS dmg	Adapt. dmg	Ratio
Cora	0.01	0%	0.33 $\pm$ 0.07	0.26 $\pm$ 0.11	0.80
Cora	0.05	0%	1.72 $\pm$ 0.35	1.32 $\pm$ 0.55	0.76
Cora	0.10	0%	3.70 $\pm$ 0.79	2.69 $\pm$ 1.18	0.72
Citeseer	0.01	0%	0.40 $\pm$ 0.08	0.34 $\pm$ 0.06	0.84
Citeseer	0.05	0%	2.14 $\pm$ 0.42	1.68 $\pm$ 0.28	0.79
Citeseer	0.10	0%	4.63 $\pm$ 0.95	3.37 $\pm$ 0.55	0.74
WikiCS	0.01	0%	0.21 $\pm$ 0.02	0.10 $\pm$ 0.10	0.47
WikiCS	0.05	0.3%	1.04 $\pm$ 0.11	0.48 $\pm$ 0.49	0.46
WikiCS	0.10	0.6%	2.10 $\pm$ 0.22	0.96 $\pm$ 0.97	0.46

are complementary: smoothing provides uniform probabilistic guarantees, AEGIS provides structurally differentiated first-order sensitivity radii for vulnerability ranking.

C. Adaptive Attack Evaluation

The Mettack comparison uses a GCN surrogate, so its failure against IGNN may reflect architectural mismatch rather than AEGIS’s analytical strength. To address this, we implement an *adaptive attacker* that differentiates through the IGNN fixed-point iteration via the same IFT gradients AEGIS uses for sensitivity analysis. This attacker optimizes perturbations directly against the IGNN loss surface using projected gradient descent [33] (PGD, 50 steps, step size $\varepsilon/10$).

The adaptive attacker breaches zero nodes at $\varepsilon = 0.01$ and fewer than 1% at $\varepsilon = 0.10$ (table III), confirming empirical validity of first-order radii. AEGIS’s SVD attack inflicts 20–50% more damage than PGD (ratio 0.46–0.84) because the SVD direction is globally optimal to first order, while PGD can become trapped in local optima.

D. Phase Transition and Scalability

Phase transition. Scaling ρ from 0.3 to 0.99 amplifies the fixed-point shift by 83 $\times$ at $\varepsilon = 0.01$, confirming the $1/(1-\kappa)$ factor from theorem 1.

Scalability. On a single NVIDIA RTX 4090 (24 GB): 1.3s/ $N=50$, 1.8s/ $N=200$ (3.7 GB), 4.4s/ $N=300$ (12.6 GB, tightness 1.000, AtkAdv 8.3 $\times$). $N=400$ exceeds 24 GB; sparse approximations would extend this.

E. Subgraph Size Ablation

Varying $N \in \{30, 50, 100, 200\}$ on Cora (10 seeds): tightness is stable at 1.01–1.02 for $N \leq 100$ and degrades to 1.031 at $N = 200$ (higher ρ). Coverage is 82–89% across all sizes. We use $N = 50$ (tightness 1.013 $\pm$ 0.003, 1.3s) as the default: 66 $\times$ faster than $N = 200$ with negligible quality loss.

F. Sensitivity and Convergence

Hyperparameters. Tightness is stable across $d \in \{16, 32, 64, 128\}$ (1.012–1.013). Varying $c \in \{0.5, 0.7, 0.9, 0.95\}$ reveals the predicted accuracy-robustness tradeoff: $c = 0.5$ gives $r_v = 0.090$ (72.1% acc); $c = 0.9$ gives $r_v = 0.051$ (80.6% acc).

TABLE IV: S_c framework applied to 7 GNN architectures (Cora, 10 seeds). Tight.: actual/predicted shift ratio. AtkAdv: AEGIS/random damage. τ : Kendall rank correlation vs. brute-force ground truth. Test Acc: full-graph accuracy. All models use edge-weighted adjacency for continuous differentiability; GAT[†] uses a non-standard edge-weighted variant (see text). Accuracy is 2–5% below published baselines due to the differentiability requirement, except GAT[†].

Model	Tight.	AtkAdv	τ	Test Acc
IGNN (equil.)	1.01 $\pm$.00	7.6 $\pm$ 0.5 $\times$	+ .32 $\pm$.04	76.9 $\pm$ 2.2%
GCN-2	1.00 $\pm$.00	4.4 $\pm$ 0.4 $\times$	+ .05 $\pm$.02	76.2 $\pm$ 0.7%
GCN-4	1.02 $\pm$.00	3.3 $\pm$ 0.2 $\times$	+ .48 $\pm$.01	78.3 $\pm$ 1.3%
GIN-2	1.00 $\pm$.00	3.8 $\pm$ 0.3 $\times$	+ .06 $\pm$.05	70.6 $\pm$ 2.4%
GAT-2 [†]	0.99 $\pm$.02	4.4 $\pm$ 1.5 $\times$	+ .36 $\pm$.13	64.2 $\pm$ 2.2%
SAGE-2	1.00 $\pm$.00	2.0 $\pm$ 0.1 $\times$	+ .49 $\pm$.10	74.4 $\pm$ 1.4%
APPNP	1.02 $\pm$.00	3.3 $\pm$ 0.2 $\times$	+ .36 $\pm$.02	80.8 $\pm$ 0.7%

Convergence. 100% convergence across all 9 datasets and 10 seeds. Pseudospectral index $\eta = 1.02$ –1.28, confirming mild non-normality and validating the κ -based bounds.

G. Defense-Informed Edge Protection

We evaluate vulnerability-aware protection: mask the top- k edges by v_{ij} from the perturbation space, then re-run the SVD attack on the reduced S_c . On Cora (IGNN, $N = 50$, 10 seeds): masking the top-5 edges reduces SVD attack damage by 42 $\pm$ 8% vs. 11 $\pm$ 6% for random 5-edge masking. At top-10: 61 $\pm$ 7% vs. 18 $\pm$ 9%. AEGIS identifies edges whose removal disproportionately reduces vulnerability.

H. Extension to Explicit GNNs

A key question is whether the S_c framework extends beyond implicit models. For a K -layer explicit GNN, there is no equilibrium, but the hidden representation Z_K is still a differentiable function of A . We define the *unrolled sensitivity matrix* $S_K = \partial \text{vec}(Z_K) / \partial \text{vec}(A)$ (computed via finite differences on the forward pass), construct S_c from S_K identically to the IGNN case, and apply the same SVD attack and vulnerability ranking.

table IV applies the S_c framework to six explicit GNN architectures on Cora. All six achieve near-perfect first-order tightness (0.99–1.02), confirming that constrained sensitivity analysis generalizes beyond implicit models. Attack advantage ranges from 2.0 $\times$ (SAGE) to 7.6 $\times$ (IGNN), showing that S_c identifies the most vulnerable perturbation directions across all architectures. Test accuracy is 2–5% below published baselines (e.g., GCN-2: 76.2% vs. 81.5% in [32]) because the edge-weighted formulation modifies the aggregation; this trade-off enables continuous differentiability for S_c computation. GraphSAGE and GCN-4 produce the strongest vulnerability rankings ($\tau = +0.49$ and $+0.48$). APPNP achieves $\tau = +0.36$ with 80.8% accuracy.

GAT requires an edge-weighted formulation where $\hat{A}$ values modulate attention weights continuously; standard GAT uses A as a binary mask (zeroing non-edges), so $\partial Z / \partial A_{ij} = 0$

for all existing edges. GATv2 [34] partially addresses this via dynamic attention, but still applies a binary mask at the neighborhood level. Our edge-weighted variant multiplies attention coefficients by $\hat{A}_{ij}$, restoring continuous differentiability. With this modification, GAT achieves tightness 0.99, attack advantage 4.4 $\times$, and $\tau = +0.36$.

Comparison with gradient attribution. Against brute-force edge-removal ground truth, AEGIS rankings achieve positive τ (IGNN: +0.32; GCN-2: +0.04; GIN-2: +0.06) while gradient-based edge attribution [35], [36] is consistently negative (−0.07, −0.17, −0.21 respectively). Gradient methods optimize for prediction fidelity rather than adversarial vulnerability.

Proposition 4 bound tightness. The ratio (Prop. 4 bound)/ $\sigma_1(S_K)$ ranges from 1.4 $\times$ (SAGE-2) to 5.9 $\times$ (APPNP), with 2-layer models achieving tighter ratios (GCN-2: 1.8 $\times$, GIN-2: 2.1 $\times$) than deeper ones (GCN-4: 4.2 $\times$, GAT-2: 3.7 $\times$). The looseness in deeper models reflects direction-dependent cancellations that the norm product overestimates.

Degree-weighted S_c . Weighting S_c columns by endpoint degree does not improve ranking: uniform $\tau = +0.042$, degree-weighted +0.037, inverse-degree −0.008 (GCN-2, 10 seeds). High-degree nodes already contribute larger sensitivity columns via aggregation.

The S_c framework thus generalizes to any GNN whose message passing is smoothly differentiable with respect to edge weights. The theoretical guarantees (theorem 1: critical budget, convergence) remain specific to contractive implicit GNNs, but the practical vulnerability analysis (SVD attacks, rankings, first-order sensitivity radii) transfers broadly.

VI. CASE STUDY: POWER FLOW CONTINGENCY

The experiments above validate AEGIS on standard graph benchmarks. A natural question is whether the vulnerability spectrum captures meaningful structure beyond these benchmarks. We demonstrate on a fundamentally different domain, AC power flow, where the adversarial vulnerability spectrum recovers a classical engineering problem: N-1 contingency analysis. Recent work has applied GNNs to power flow computation [37], [38] and contingency screening [3]; AEGIS approaches the same problem from the vulnerability-analysis perspective.

A. Background: N-1 Contingency

Power grid operators must ensure that no single transmission line failure causes system instability. The standard N-1 contingency assessment removes each line in turn, re-solves the power flow equations, and ranks lines by the severity of the resulting voltage and flow deviations [39]. This brute-force procedure scales as $O(|E|)$ full power flow solves, which is expensive for large grids and infeasible for real-time operation.

We observe that AEGIS’s per-edge vulnerability $v_{ij} = \|[S_c]_{:,k}\|_2$ computes an analogous quantity via a single IFT analysis: edge perturbation maps to line trips, the vulnerability spectrum maps to contingency severity, and critical budget

TABLE V: AEGIS on IEEE power flow benchmarks (10 seeds). Model quality (per-unit RMSE): $|V|$ = voltage magnitude, θ = angle, ΔS = power-balance residual per bus. τ : Kendall rank correlation vs. brute-force N-1 contingency. P@10: top-10 precision.

Case	N	$ E $	Model (p.u.)			AEGIS	
			$ V $	θ	ΔS	τ	P@10
case14	14	20	.007	.020	.106	$+.42_{\pm.19}$	$.74_{\pm.12}$
case30	30	41	.014	.024	.042	$+.37_{\pm.17}$	$.68_{\pm.06}$
case57	57	78	.033	.059	.033	$+.67_{\pm.09}$	$.66_{\pm.13}$
case118	118	179	.014	.076	.042	$+.62_{\pm.11}$	$.81_{\pm.10}$

$\varepsilon_{\text{crit}}$ maps to stability margin. The correspondence is approximate (continuous first-order vs. discrete removal), analogous to how PTDF/LODF [39] linearize the effect of line trips.

B. Experimental Setup

We train ContractiveGCN-PF, an IGNN variant for AC power flow, on four standard IEEE test cases [30]. Training data is generated via PandaPower’s [31] full AC Newton-Raphson solver (2,000 load samples per case, uniformly sampled at 70–130% of nominal load). The model takes binary adjacency from the admittance matrix and 5 bus features: bus type indicators (slack, PV), active and reactive power injections (P, Q), and voltage magnitude $|V|$. Unlike graph benchmarks, power flow grids are small enough (14–118 buses) to analyze without subgraph extraction.

Limitation. Training data covers uniform load scaling (2,000 samples) but not seasonal or generator-outage variation.

C. Results

table V reports constrained tightness and N-1 ranking correlation (Kendall τ between AEGIS vulnerability spectrum and brute-force contingency ranking) across 10 seeds.

Key observations. First, tightness transfers to the power flow domain (≈ 1.00), confirming that the first-order prediction is not domain-specific. Second, ranking correlation ($\tau = 0.37$ – 0.67) is moderate but the operationally relevant P@10 is 0.66–0.81, correctly identifying 7–8 of the top 10 critical lines.

Implicit physics from equilibrium structure. The model is trained with voltage MSE alone — no power-balance penalty. Yet the per-bus residual $\Delta S = 0.03$ – 0.11 p.u. (table V), meaning predicted voltages approximately satisfy Kirchhoff’s laws without explicit enforcement. This arises from the DEQ architecture: at equilibrium, $Z^* = F(Z^*, A)$ forces the hidden state to be self-consistent with the grid topology in A . Perturbations that disrupt this self-consistency (identified by AEGIS) are precisely those that break the power balance, connecting structural vulnerability directly to physical contingency.

Baseline comparison. LODF (industry-standard DC screening [39]) achieves $\tau = 0.44$ – 0.58 ; AEGIS outperforms on larger grids ($\tau = 0.62$ – 0.67 on case57/118) by capturing nonlinear AC effects. LODF requires line reactances; AEGIS requires only the trained model.

Practical positioning. For case118, AEGIS ranks in 2.3s (single IFT) vs. 179 full AC solves for brute-force N-1, achieving P@10 = 0.81. Binary adjacency outperforms admittance-weighted (P@10 = 0.81 vs. 0.27) because N-1 contingency is a discrete event better modeled by uniform sensitivity.

Operational caveat. AEGIS is a screening layer, not a standalone contingency tool ($\tau = 0.37$ – 0.67 is insufficient for direct operational use).

VII. RELATED WORK

AEGIS addresses adversarial structural vulnerability analysis for GNNs through first-order sensitivity. We position it against five research threads: structural attacks, certified robustness, empirical defenses, implicit network theory, and sensitivity analysis.

Adversarial attacks on graph structure. Zügner et al. [1] demonstrated GNN vulnerability to structural perturbation with Nettack; Mettack [4] extended this to global poisoning via meta-gradients through a GCN [32] surrogate. Dai et al. [2] framed attacks as reinforcement learning, Bojchevski and Günnemann [40] proposed global poisoning via spectral shifts, and Xu et al. [41] formulated topology attack as bilevel optimization. Black-box variants include GF-Attack [42] via graph signal processing and influence-function approaches [43]. Wu et al. [44] and Jin et al. [45] survey the landscape. These methods use surrogate gradients or heuristic search to find adversarial edges. AEGIS takes a fundamentally different approach: it computes the closed-form optimal attack direction via SVD of the constrained sensitivity matrix S_c , which is architecture-agnostic and applies to any differentiable GNN.

Certified robustness for GNNs. Bojchevski and Günnemann [8] pioneered certifiable robustness via convex relaxations for GCN-class models. Zügner and Günnemann [7] extended this with interval bound propagation. Randomized smoothing [5] provides probabilistic certificates for any classifier; Scholten et al. refined this with message-interception [6] and hierarchical [46] smoothing. Jia et al. [47] certified community detection via smoothing, Wang et al. [9] proposed message-passing invariance certificates, and Geisler et al. [48] scaled certified defenses to large graphs. Li et al. [49] proposed AGNNCert with deterministic guarantees. These methods provide robustness guarantees but do not identify *which* edges or nodes are vulnerable. AEGIS complements them by providing structurally differentiated per-node radii and per-edge vulnerability rankings that reveal the attack surface geometry (section V-B).

GNN defense. Empirical defenses improve robustness without formal guarantees: Robust GCN [10] absorbs noise via Gaussian hidden layers, Pro-GNN [12] jointly learns clean structure, GNNGuard [11] reweights edges during aggregation, PA-GNN [50] transfers robustness across domains, PTDNet [13] learns denoised topologies, and soft-median aggregation [51] replaces mean pooling. Entezari et al. [52] observed that adversarial perturbations primarily affect high-rank spectral components. AEGIS is an analytical tool rather than a defense, but its per-edge vulnerability spectrum could

inform defense design by identifying which edges to protect, and the SVD structure echoes the spectral observations of Entezari et al. from the sensitivity perspective.

Implicit networks and equilibrium analysis. The DEQ framework [15] introduced fixed-point models; extensions include multiscale DEQs [16], Jacobian regularization [20], phantom gradients [53], and Jacobian-free backpropagation [19]. IGNN [14] enforces unique equilibria via spectral normalization [17]; EIGNN [54] improves efficiency. For robustness, monotone operator networks [55] provide built-in Lipschitz bounds, Revay et al. [24] use LMIs, Pabbaraju et al. [56] estimate tighter constants, and El Ghaoui et al. [23] analyze well-posedness. All address input sensitivity $\|\partial z^*/\partial x\|$, not structural sensitivity $\|\partial Z/\partial A\|$. AEGIS addresses structural sensitivity for both implicit models (via the resolvent $(I - J_z)^{-1}J_A$) and explicit models (via the unrolled Jacobian S_K), providing stronger theoretical guarantees for the implicit case (theorem 1).

Sensitivity analysis and explainability. AEGIS draws on classical matrix perturbation theory [25], [21]. Novak et al. [57] connected Jacobian sensitivity to generalization; Liu et al. [58] studied GNN stability without per-node radii. GNN explanation methods produce per-edge importance scores: GNNExplainer [36] learns soft masks, PGExplainer [59] trains parametric edge-mask generators, and gradient-based methods [35] use input gradients. These optimize for *prediction fidelity* (which edges matter for the output) rather than *adversarial vulnerability* (which perturbation maximizes hidden-state shift). AEGIS’s S_c is perturbation-space optimal by construction (SVD), provides quantitative sensitivity radii, and identifies the globally optimal attack direction.

GNNs for power systems. Our case study (section VI) connects to GNN-based power flow solvers [37], [38] and GNN contingency screening [3]. AEGIS shows that structural sensitivity independently recovers contingency rankings without domain-specific linearization.

Positioning. AEGIS provides structural vulnerability analysis via the constrained sensitivity matrix S_c : for any differentiable GNN, S_c yields SVD-optimal attacks, per-edge rankings, and first-order sensitivity radii (theorem 4). For contractive implicit GNNs, the IFT additionally provides a critical budget $\varepsilon_{\text{crit}}$ and convergence guarantees (theorem 1) that explicit architectures lack. The distinction is important: for explicit GNNs, AEGIS is a computational tool for vulnerability analysis; for implicit GNNs, it additionally offers formal regime characterization. Validated across 7 architectures (section V-H), AEGIS offers a combination of generality and structural differentiation that no existing method provides, with the strongest theoretical backing in the implicit case.

VIII. CONCLUSION

AEGIS provides a principled pre-deployment diagnostic for GNNs on safety-critical graphs: from a single S_c computation, practitioners obtain the optimal attack direction, per-edge vulnerability rankings, and per-node tolerance budgets. Predictions are tight across 7 architectures and 9 datasets

(tightness 1.00 ± 0.01 at $\varepsilon = 0.01$), and the vulnerability spectrum transfers to power grid contingency analysis ($P@10 = 0.66\text{--}0.81$). The per-edge ranking directly informs defense design (section V-G).

Limitations. (1) First-order sensitivity radii are locally tight but not global certificates without second-order bounds. (2) Continuous edge-weight perturbations; discrete insertions/deletions are outside the formal guarantee. (3) Dense Jacobian limits subgraph size to $N \approx 300$. (4) Standard GAT requires edge-weighted modification for continuous sensitivity. (5) Power flow $\tau = 0.37\text{--}0.67$ requires verification for operational use [39].

Future work. Sparse solvers for larger subgraphs; second-order bounds for formal certificates; discrete perturbation models; integration with defense mechanisms [12], [11] for vulnerability-aware training.

REFERENCES

- [1] D. Zügner, A. Akbarnejad, and S. Günnemann, “Adversarial attacks on neural networks for graph data,” in *Proceedings of the 24th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, 2018, pp. 2847–2856.
- [2] H. Dai, H. Li, T. Tian, X. Huang, L. Wang, J. Zhu, and L. Song, “Adversarial attack on graph structured data,” in *International Conference on Machine Learning*, 2018, pp. 1115–1124.
- [3] A. M. Nakiganda, C. Weeraddana, N. Popli, and T. van der Hagen, “Graph neural networks for fast contingency analysis of power systems,” *arXiv preprint arXiv:2310.04213*, 2023.
- [4] D. Zügner and S. Günnemann, “Adversarial attacks on graph neural networks via meta learning,” in *International Conference on Learning Representations*, 2019.
- [5] A. Bojchevski, J. Klicpera, and S. Günnemann, “Efficient robustness certificates for discrete data: Sparsity-aware randomized smoothing for graphs, images and more,” in *International Conference on Machine Learning*, 2020, pp. 1003–1013.
- [6] Y. Scholten, J. Schuchardt, A. Bojchevski, and S. Günnemann, “Randomized message-interception smoothing: Gray-box certificates for graph neural networks,” in *Advances in Neural Information Processing Systems*, vol. 35, 2022.
- [7] D. Zügner and S. Günnemann, “Certifiable robustness and robust training for graph convolutional networks,” in *Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, 2019, pp. 246–256.
- [8] A. Bojchevski and S. Günnemann, “Certifiable robustness to graph perturbations,” in *Advances in Neural Information Processing Systems*, vol. 32, 2019.
- [9] B. Wang, J. Jia, X. Cao, and N. Z. Gong, “Certified robustness of graph neural networks against adversarial structural perturbation,” in *Proceedings of the 27th ACM SIGKDD Conference on Knowledge Discovery & Data Mining*, 2021, pp. 1645–1653.
- [10] D. Zhu, Z. Zhang, P. Cui, and W. Zhu, “Robust graph convolutional networks against adversarial attacks,” in *Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, 2019, pp. 1399–1407.
- [11] X. Zhang and M. Zitnik, “GNNGuard: Defending graph neural networks against adversarial attacks,” in *Advances in Neural Information Processing Systems*, vol. 33, 2020, pp. 9263–9275.
- [12] W. Jin, Y. Ma, X. Liu, X. Tang, S. Wang, and J. Tang, “Graph structure learning for robust graph neural networks,” in *Proceedings of the 26th ACM SIGKDD Conference on Knowledge Discovery & Data Mining*, 2020, pp. 66–74.
- [13] D. Luo, W. Cheng, W. Yu, B. Zong, J. Ni, H. Chen, and X. Zhang, “Learning to drop: Robust graph neural network via topological denoising,” in *Proceedings of the 14th ACM International Conference on Web Search and Data Mining*, 2021, pp. 779–787.
- [14] F. Gu, H. Chang, W. Zhu, S. Sojoudi, and L. El Ghaoui, “Implicit graph neural networks,” in *Advances in Neural Information Processing Systems*, vol. 33, 2020, pp. 11 984–11 995.

- [15] S. Bai, J. Z. Kolter, and V. Koltun, "Deep equilibrium models," in *Advances in Neural Information Processing Systems*, vol. 32, 2019.
- [16] S. Bai, V. Koltun, and J. Z. Kolter, "Multiscale deep equilibrium models," in *Advances in Neural Information Processing Systems*, vol. 33, 2020, pp. 5238–5250.
- [17] T. Miyato, T. Kataoka, M. Koyama, and Y. Yoshida, "Spectral normalization for generative adversarial networks," in *International Conference on Learning Representations*, 2018.
- [18] S. Gould, R. Hartley, and D. Campbell, "Deep declarative networks," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 44, no. 8, pp. 3988–4004, 2021.
- [19] S. W. Fung, H. Heaton, Q. Li, D. McKenzie, S. Osher, and W. Yin, "JFB: Jacobian-free backpropagation for implicit networks," in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 36, no. 6, 2022, pp. 6648–6656.
- [20] S. Bai, V. Koltun, and J. Z. Kolter, "Stabilizing equilibrium models by Jacobian regularization," in *International Conference on Machine Learning*, 2021, pp. 554–565.
- [21] L. N. Trefethen and M. Embree, *Spectra and Pseudospectra: The Behavior of Nonnormal Matrices and Operators*. Princeton University Press, 2005.
- [22] J. Lorraine, P. Vicol, and D. Duvenaud, "Optimizing millions of hyperparameters by implicit differentiation," in *International Conference on Artificial Intelligence and Statistics*, 2020, pp. 1540–1552.
- [23] L. El Ghaoui, F. Gu, B. Travacca, A. Askari, and A. Tsai, "Implicit deep learning," *SIAM Journal on Mathematics of Data Science*, vol. 3, no. 3, pp. 930–958, 2021.
- [24] M. Revay, R. Wang, and I. R. Manchester, "Lipschitz bounded equilibrium networks," in *arXiv preprint arXiv:2010.01732*, 2020.
- [25] G. W. Stewart and J.-g. Sun, *Matrix Perturbation Theory*. Academic Press, 1990.
- [26] A. Paszke, S. Gross, F. Massa, A. Lerner, J. Bradbury, G. Chanan, T. Killeen, Z. Lin, N. Gimelshein, L. Antiga *et al.*, "PyTorch: An imperative style, high-performance deep learning library," in *Advances in Neural Information Processing Systems*, vol. 32, 2019.
- [27] P. Sen, G. Namata, M. Bilgic, L. Getoor, B. Galligher, and T. Eliassirad, "Collective classification in network data," *AI Magazine*, vol. 29, no. 3, pp. 93–106, 2008.
- [28] O. Shchur, M. Mumme, A. Bojchevski, and S. Günnemann, "Pitfalls of graph neural network evaluation," in *Relational Representation Learning Workshop, NeurIPS*, 2018.
- [29] P. Mernyei and M. Caron, "Wiki-CS: A Wikipedia-based benchmark for graph neural networks," in *arXiv preprint arXiv:2007.02901*, 2020.
- [30] IEEE PES, "IEEE PES test feeder working group: Power system test cases," *IEEE Power and Energy Society*, 2024, <https://electricgrids.engr.tamu.edu>.
- [31] L. Thurner, A. Scheidler, F. Schäfer, J.-H. Menke, J. Dolichon, F. Meier, S. Meinecke, and M. Braun, "pandapower – an open-source Python tool for convenient modeling, analysis, and optimization of electric power systems," *IEEE Transactions on Power Systems*, vol. 33, no. 6, pp. 6510–6521, 2018.
- [32] T. N. Kipf and M. Welling, "Semi-supervised classification with graph convolutional networks," in *International Conference on Learning Representations*, 2017.
- [33] A. Madry, A. Makelov, L. Schmidt, D. Tsipras, and A. Vladu, "Towards deep learning models resistant to adversarial attacks," in *International Conference on Learning Representations*, 2018.
- [34] S. Brody, U. Alon, and E. Yahav, "How attentive are graph attention networks?" in *International Conference on Learning Representations*, 2022.
- [35] P. E. Pope, S. Kolouri, M. Rostami, C. E. Martin, and H. Hoffmann, "Explainability methods for graph convolutional neural networks," in *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2019, pp. 10772–10781.
- [36] R. Ying, D. Bourgeois, J. You, M. Zitnik, and J. Leskovec, "GN-Explainer: Generating explanations for graph neural networks," in *Advances in Neural Information Processing Systems*, vol. 32, 2019.
- [37] B. Donon, B. Donnot, I. Guyon, and A. Marot, "Graph neural solver for power systems," in *International Joint Conference on Neural Networks*, 2019, pp. 1–8.
- [38] V. Bolz, J. Rueß, and A. Zell, "Power flow approximation based on graph convolutional networks," in *18th IEEE International Conference on Machine Learning and Applications*, 2019, pp. 1679–1686.
- [39] A. J. Wood, B. F. Wollenberg, and G. B. Sheblé, *Power Generation, Operation, and Control*, 3rd ed. John Wiley & Sons, 2014.
- [40] A. Bojchevski and S. Günnemann, "Adversarial attacks on node embeddings via graph poisoning," in *International Conference on Machine Learning*, 2019, pp. 695–704.
- [41] K. Xu, H. Chen, S. Liu, P.-Y. Chen, T.-W. Weng, M. Hong, and X. Lin, "Topology attack and defense for graph neural networks: An optimization perspective," in *International Joint Conference on Artificial Intelligence*, 2019, pp. 3961–3967.
- [42] H. Chang, Y. Rong, T. Xu, W. Huang, H. Zhang, P. Cui, W. Zhu, and J. Huang, "A restricted black-box adversarial framework towards attacking graph embedding models," in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 34, no. 04, 2020, pp. 3389–3396.
- [43] J. Ma, S. Ding, and Q. Mei, "Towards more practical adversarial attacks on graph neural networks," in *Advances in Neural Information Processing Systems*, vol. 33, 2020, pp. 4756–4766.
- [44] H. Wu, C. Wang, Y. Tyshetskiy, A. Docherty, K. Lu, and L. Zhu, "Adversarial examples on graph data: Deep insights into attack and defense," in *International Joint Conference on Artificial Intelligence*, 2019, pp. 4816–4823.
- [45] W. Jin, Y. Li, H. Xu, Y. Wang, S. Ji, C. Aggarwal, and J. Tang, "Adversarial attacks and defenses on graphs," *ACM SIGKDD Explorations Newsletter*, vol. 22, no. 2, pp. 19–34, 2021.
- [46] Y. Scholten, J. Schuchardt, A. Bojchevski, and S. Günnemann, "Hierarchical randomized smoothing," in *Advances in Neural Information Processing Systems*, vol. 36, 2023.
- [47] J. Jia, B. Wang, X. Cao, and N. Z. Gong, "Certified robustness of community detection against adversarial structural perturbation via randomized smoothing," in *Proceedings of The Web Conference*, 2020, pp. 2718–2724.
- [48] S. Geisler, T. Schmidt, H. Sirin, D. Zügner, A. Bojchevski, and S. Günnemann, "Robustness of graph neural networks at scale," in *Advances in Neural Information Processing Systems*, vol. 34, 2021, pp. 7637–7649.
- [49] J. Li, G. Tao, D. Han, and X. Zhang, "AGNNCert: Defending graph neural networks against arbitrary perturbations with certified robustness," in *USENIX Security Symposium*, 2025.
- [50] X. Tang, Y. Li, Y. Sun, H. Yao, P. Mitra, and S. Wang, "Transferring robustness for graph neural network against poisoning attacks," in *Proceedings of the 13th International Conference on Web Search and Data Mining*, 2020, pp. 600–608.
- [51] S. Geisler, D. Zügner, and S. Günnemann, "Reliable graph neural networks via robust aggregation," in *Advances in Neural Information Processing Systems*, vol. 33, 2020, pp. 13 272–13 284.
- [52] N. Entezari, S. A. Al-Sayouri, A. Darvishzadeh, and E. E. Papalexakis, "All you need is low (rank): Defending against adversarial attacks on graphs," in *Proceedings of the 13th International Conference on Web Search and Data Mining*, 2020, pp. 169–177.
- [53] Z. Geng, X.-Y. Zhang, S. Bai, Y. Wang, and Z. Lin, "On training implicit models," in *Advances in Neural Information Processing Systems*, vol. 34, 2021, pp. 24 247–24 260.
- [54] J. Liu, K. Kawaguchi, B. Hooi, Y. Wang, and J. Feng, "Eignn: Efficient infinite-depth graph neural networks," in *Advances in Neural Information Processing Systems*, vol. 34, 2021, pp. 18 762–18 773.
- [55] E. Winston and J. Z. Kolter, "Monotone operator equilibrium networks," in *Advances in Neural Information Processing Systems*, vol. 33, 2020, pp. 10 718–10 728.
- [56] C. Pabbaraju, E. Winston, and J. Z. Kolter, "Estimating lipschitz constants of monotone deep equilibrium models," *arXiv preprint arXiv:2111.08931*, 2021.
- [57] R. Novak, Y. Bahri, D. A. Abolafia, J. Pennington, and J. Sohl-Dickstein, "Sensitivity and generalization in neural networks: An empirical study," in *International Conference on Learning Representations*, 2018.
- [58] S. Liu, R. Ying, H. Dong, L. Li, T. Xu, Y. Liang, J. Leskovec, and H. Peng, "Stability and generalization of graph neural networks," in *arXiv preprint*, 2021.
- [59] D. Luo, W. Cheng, D. Xu, W. Yu, B. Zong, H. Chen, and X. Zhang, "Parameterized explainer for graph neural network," in *Advances in Neural Information Processing Systems*, vol. 33, 2020, pp. 19 620–19 631.