

AEGIS: Equilibrium Structure Mining for Certifiable Graph Robustness

Redwanul Karim*, [Co-Author]*

*[Affiliation]

Email: karim.redwanul.de@gmail.com

Abstract—We introduce AEGIS, a framework that mines the equilibrium structure of Deep Equilibrium Graph Neural Networks (DEQ-GNNs) to provide exact first-order prediction of adversarial vulnerability under graph structure perturbation. Unlike explicit GNNs, DEQ-GNNs converge to a fixed point $z^* = F(z^*, A)$, enabling construction of a structural sensitivity matrix $S = (I - J_z)^{-1} J_A$ via the implicit function theorem. We prove that structural perturbations exhibit a sharp phase transition (Theorem 1): below a critical budget $\varepsilon_{\text{crit}} = (1 - \rho) / \|W\|_2$, fixed-point shifts are bounded and certifiable; above it, contractivity is lost and certificates become void. Empirically, our constrained sensitivity analysis achieves tightness 1.00 ± 0.01 across 9 datasets (5 graph benchmarks + 4 IEEE power grids, 10 seeds each), meaning the predicted adversarial shift matches the actual shift to within 1%. AEGIS attacks outperform Mettack in 30/30 comparisons, and its deterministic per-node certificates are $2\text{--}5\times$ larger than randomized smoothing at equal coverage. As a case study, the vulnerability spectrum naturally recovers power grid N-1 contingency rankings ($\tau = 0.42\text{--}0.67$). Code is available at [URL].

Index Terms—Deep equilibrium models, graph neural networks, adversarial robustness, certified robustness, structural perturbation, implicit function theorem

I. INTRODUCTION

Graph neural networks (GNNs) are vulnerable to adversarial structural perturbation: small changes to the graph topology—adding or removing a few edges—can drastically alter predictions [1], [2]. Defending against such attacks requires understanding *which perturbations matter* and *by how much*. Current approaches either provide loose certified bounds via randomized smoothing [3] or rely on heuristic attacks without guarantees [2].

We observe that a specific family of GNNs—Deep Equilibrium GNNs (DEQ-GNNs) [4], [5]—possesses a structural property that enables a qualitatively different kind of analysis. These models iterate a contractive operator F_θ until convergence to a fixed point $z^* = F_\theta(z^*, A)$, rather than propagating through a fixed number of layers. This fixed-point structure, combined with the implicit function theorem (IFT), yields the *structural sensitivity matrix*

$$S = (I - J_z)^{-1} J_A, \quad (1)$$

where $J_z = \partial F / \partial z|_{z^*}$ and $J_A = \partial F / \partial \text{vec}(A)|_{z^*}$. The matrix S maps structural perturbations to fixed-point shifts: $\Delta z^* \approx S \cdot \text{vec}(\delta A)$.

The central finding of this paper is that, under realistic (symmetric, edge-only) perturbation constraints, S predicts adversarial vulnerability *exactly* to first order. Across 9 datasets

spanning 4 domains and 10 random seeds, the ratio of actual to predicted adversarial shift is 1.00 ± 0.01 . No other GNN robustness method achieves this level of prediction accuracy.

Contributions. We make three contributions:

- 1) **Phase transition theorem.** We prove that structural perturbations to contractive DEQ-GNNs exhibit three regimes around a critical budget $\varepsilon_{\text{crit}} = (1 - \rho) / \|W\|_2$: subcritical (bounded, certifiable), critical (divergent), and supercritical (certificates void). This is the first characterization of this phase transition for DEQ-GNNs (section III).
- 2) **AEGIS framework.** We develop a domain-agnostic framework for mining DEQ equilibrium structure, providing optimal structural attacks via SVD of the constrained sensitivity matrix and deterministic per-node robustness certificates. AEGIS attacks outperform Mettack [2] in 30/30 comparisons across 3 datasets and 10 seeds. Its certificates are $2\text{--}5\times$ larger than randomized smoothing [3] at equal coverage (sections IV and V).
- 3) **Cross-domain validation.** We validate on 9 datasets across 4 domains—citation networks (Cora, Citeseer, Pubmed), e-commerce (Amazon Photo), encyclopedia (WikiCS), and power flow (IEEE case14/30/57/118)—demonstrating that constrained tightness ≈ 1.00 holds universally. As a case study, the vulnerability spectrum recovers power grid N-1 contingency rankings (section V).

II. BACKGROUND

A. Deep Equilibrium Graph Neural Networks

A DEQ-GNN [4] defines a weight-tied operator $F_\theta : \mathbb{R}^{N \times d} \times \mathbb{R}^{N \times N} \rightarrow \mathbb{R}^{N \times d}$ and iterates it to convergence:

$$Z_{k+1} = F_\theta(Z_k, A), \quad Z_0 = 0. \quad (2)$$

When F_θ is contractive (spectral radius $\rho(J_z) < 1$), the Banach fixed-point theorem guarantees a unique equilibrium $z^* = F_\theta(z^*, A)$. The Implicit Graph Neural Network (IGNN) [5] instantiates this with $F(Z, A) = \sigma(\hat{A}ZW^\top + X_{\text{proj}})$, where $\hat{A} = D^{-1/2}(A + I)D^{-1/2}$ is the normalized adjacency and W is spectral-norm constrained to ensure $\rho < 1$.

B. Implicit Function Theorem for DEQ Sensitivity

At the fixed point, the IFT provides the sensitivity of z^* to any parameter p :

$$\frac{\partial z^*}{\partial p} = (I - J_z)^{-1} \frac{\partial F}{\partial p} \Big|_{z^*}, \quad (3)$$

where $J_z = \partial F / \partial z|_{z^*}$ is the state Jacobian. When $\rho(J_z) < 1$, the resolvent $(I - J_z)^{-1}$ exists and is bounded by $\|(I - J_z)^{-1}\|_2 \leq 1/(1 - \rho)$ for normal J_z . This has been used for hyperparameter optimization in implicit layers [6] and general implicit differentiation [7], but not for adversarial robustness analysis of graph structure.

C. Adversarial Structural Perturbation

We consider an attacker who perturbs the adjacency matrix: $A' = A + \delta A$ with $\|\delta A\|_F \leq \varepsilon$. For undirected graphs, δA must be symmetric. In graph terms, this corresponds to modifying edge weights or adding/removing edges. Unlike input-feature perturbation (studied for DEQs in [8], [9]), structural perturbation changes the *operator* itself, not just its input—the fixed point depends on A through both J_z and J_A .

III. ADVERSARIAL EQUILIBRIUM THEORY

We present our main theoretical result: a phase transition in adversarial vulnerability, supported by two propositions providing the computational tools.

[Phase Transition in Adversarial Vulnerability] Let $F_\theta(Z, A) = \sigma(AZW^\top + X_{\text{proj}})$ be an IGNN-class operator with $\|W\|_2 \leq c$ and spectral radius $\rho = \rho(J_z) < 1$. Define

$$\varepsilon_{\text{crit}} = \frac{1 - \rho}{\|W\|_2}. \quad (4)$$

Then structural perturbations δA exhibit three regimes:

(a) **Subcritical** ($\|\delta A\|_F < \varepsilon_{\text{crit}}$): The perturbed operator remains contractive with a unique fixed point satisfying

$$\|\Delta z^*\| \leq \sigma_1(S) \cdot \|\delta A\|_F + O(\|\delta A\|_F^2), \quad (5)$$

where $S = (I - J_z)^{-1} J_A$ and $\sigma_1(S) \leq \|J_A\|_{\text{op}} / (1 - \rho)$.

(b) **Critical** ($\|\delta A\|_F \rightarrow \varepsilon_{\text{crit}}$): The bound diverges as $\Theta(1/(\varepsilon_{\text{crit}} - \|\delta A\|_F))$ because $\rho(A + \delta A) \rightarrow 1$.

(c) **Supercritical** ($\|\delta A\|_F > \varepsilon_{\text{crit}}$): Contractivity may be lost ($\rho \geq 1$), the fixed point is no longer guaranteed unique, and all certificates from part (a) become void.

Part (a). By the IFT applied to $z^* = F(z^*, A)$:

$$\Delta z^* = (I - J_z)^{-1} J_A \cdot \text{vec}(\delta A) + O(\|\delta A\|^2). \quad (6)$$

Taking norms gives (5). For contractivity preservation: $\rho(A + \delta A) \leq \|A + \delta A\|_2 \cdot \|W\|_2 \leq (\|A\|_2 + \|\delta A\|_2) \cdot \|W\|_2$. Since $\rho = \|A\|_2 \cdot \|W\|_2$ in the worst case and $\|\cdot\|_2 \leq \|\cdot\|_F$, the condition $\|\delta A\|_F < (1 - \rho) / \|W\|_2$ ensures $\rho(A + \delta A) < 1$.

Part (b). As $\|\delta A\|_F \rightarrow \varepsilon_{\text{crit}}$, $\rho(A + \delta A) \rightarrow 1$ and $\|(I - J_z(A + \delta A))^{-1}\|_2 \geq 1/(1 - \rho(A + \delta A)) \rightarrow \infty$.

Part (c). When $\rho \geq 1$, the contraction mapping theorem no longer applies. The fixed-point iteration may converge to a different equilibrium, oscillate, or diverge.

[Optimal First-Order Structural Attack] The perturbation δA^* maximizing $\|\Delta z^*\|$ to first order subject to $\|\delta A\|_F \leq \varepsilon$ is $\delta A^* = \varepsilon \cdot \text{reshape}(v_1, N \times N)$, where v_1 is the leading right singular vector of S . The maximum first-order shift is $\varepsilon \cdot \sigma_1(S)$.

Constrained sensitivity. Real graph perturbations are symmetric ($\delta A = \delta A^\top$) and edge-only (only existing edges can be modified). We construct the *constrained sensitivity matrix* $S_c \in \mathbb{R}^{D \times |E|}$ with column $S_{:,iN+j} + S_{:,jN+i}$ for each edge $(i, j) \in E$. The constrained-optimal attack uses $\sigma_1(S_c)$, yielding tightness ≈ 1.00 empirically (section V-A).

[Per-Node Robust Radius] For node v with classification margin $m_v = f_{y_v}(z_v^*) - \max_{c \neq y_v} f_c(z_v^*) > 0$, the certified robust radius is

$$r_v = \frac{m_v}{\|\partial f / \partial z_v^*\| \cdot \|S_v\|}, \quad (7)$$

where S_v denotes the block-rows of S for node v . Any $\|\delta A\|_F < r_v$ preserves the classification of node v . These certificates are deterministic, unlike randomized smoothing [3] which holds with probability $1 - \alpha$.

IV. AEGIS FRAMEWORK

AEGIS packages the theory into a domain-agnostic Python API. Given any contractive DEQ-GNN with operator F , fixed point z^* , and context (adjacency, features), the `IEMiner` class provides:

- 1) **Vulnerability analysis:** compute S , $\sigma_1(S_c)$, and the per-edge vulnerability spectrum $v_{ij} = \|S_{:,iN+j} + S_{:,jN+i}\|$.
- 2) **Certified robustness:** per-node radii r_v via section III, critical budget $\varepsilon_{\text{crit}}$ via section III.
- 3) **Contractivity verification:** spectral radius ρ via power iteration, non-normality index $\eta = \|(I - J_z)^{-1}\|_2 \cdot (1 - \rho)$.

Subgraph-based analysis. The framework operates on BFS ego-subgraphs (typically 50–200 nodes) extracted from the full graph. The model is trained on the full graph; only the adversarial analysis is local. This suffices because per-node certificates and edge vulnerability are inherently local quantities—a node’s certified radius depends on its own margin and the sensitivity of its neighborhood, not distant parts of the graph.

Computational cost. The dominant cost is computing $J_z \in \mathbb{R}^{D \times D}$ and $J_A \in \mathbb{R}^{D \times N^2}$ via row-by-row backward passes, followed by a linear solve for S . Total cost is $O(D \cdot N^2 + D^3)$ where $D = N \cdot d$. On a single GPU, this ranges from 0.5s ($N = 20$) to 8.1s ($N = 200$); see section V-E.

V. EXPERIMENTS

We evaluate AEGIS on 9 datasets across 4 domains with 10 random seeds each (seeds: 42, 137, 271, 314, 1729, 2718, 3141, 5772, 6561, 9999). All graph benchmark experiments use IGNN [5] with hidden dimension 64, spectral-norm constrained weights, and early stopping on validation accuracy over 200 epochs. Power flow experiments use ContractiveGCN-PF (an IGNN variant for AC power flow).

TABLE I: Cross-domain adversarial analysis (mean $\pm$ std, 10 seeds). Constrained tightness ≈ 1.00 across all domains.

Dataset	ρ	Tight.	AtkAdv	$\varepsilon_{\text{crit}}$	Cert%
Cora	.33 $\pm$.14	1.01 $\pm$.01	3.6 $\pm$.6x	.66 $\pm$.15	81 $\pm$ 9%
Citeseer	.59 $\pm$.02	1.01 $\pm$.01	4.1 $\pm$.5x	.41 $\pm$.02	83 $\pm$ 4%
Pubmed	.41 $\pm$.11	1.01 $\pm$.01	3.2 $\pm$.4x	.59 $\pm$.11	74 $\pm$ 23%
Amazon	.14 $\pm$.02	1.00 $\pm$.00	3.8 $\pm$.2x	.86 $\pm$.02	92 $\pm$ 4%
WikiCS	.34 $\pm$.04	1.00 $\pm$.01	3.8 $\pm$.4x	.66 $\pm$.04	70 $\pm$ 3%

TABLE II: IFT vs. Mettack attack at $k = 1$ edge removal (10 seeds). IFT wins 30/30 comparisons. τ : Kendall correlation with brute-force ground truth.

Dataset	IFT τ	Met τ	IFT dmg	Met dmg
Cora	+.28 $\pm$.07	+.20 $\pm$.11	.55 $\pm$.27	.16 $\pm$.09
Citeseer	+.38 $\pm$.12	+.16 $\pm$.13	1.1 $\pm$ 1.5	.11 $\pm$.14
WikiCS	+.14 $\pm$.04	+.52 $\pm$.02	.05 $\pm$.01	.01 $\pm$.00

Datasets. Graph benchmarks: Cora (2,708 nodes), Citeseer (3,327), Pubmed (19,717), Amazon Photo (7,650), WikiCS (11,701). Power flow: IEEE case14 (14 buses), case30 (30), case57 (57), case118 (118), generated via PandaPower with 2,000 samples each.

A. Cross-Domain Adversarial Analysis

table I reports the constrained tightness ratio (actual shift / predicted shift at $\varepsilon = 0.01$), attack advantage over random perturbation, and per-node certificate coverage across 50-node BFS ego-subgraphs.

Tightness is 1.00 ± 0.01 across all datasets, confirming that the constrained IFT prediction is essentially exact under symmetric edge perturbations. The attack advantage (3.2–4.1 $\times$ over random) demonstrates that AEGIS identifies the most vulnerable perturbation directions. Certificate coverage ranges from 70% (WikiCS) to 92% (Amazon Photo), with higher coverage on datasets where the model achieves higher accuracy.

B. Attack Comparison: AEGIS vs. Mettack

We compare AEGIS’s IFT-based structural attack against Mettack [2] (Meta-Self variant with 2-layer GCN surrogate). Both methods rank edges for removal; damage is measured by IGNN fixed-point shift $\|\Delta z^*\|$ after reconvergence.

AEGIS inflicts 3–10 $\times$ more damage than Mettack at every seed and budget level (30/30 wins at $k = 1, \dots, 5$). Mettack’s GCN surrogate occasionally ranks edges closer to ground truth (WikiCS $\tau = +0.52$), but the actual damage is 5 $\times$ lower because the surrogate decision boundary diverges from the equilibrium structure. Direct equilibrium analysis beats surrogate transfer.

C. Certificate Comparison: Deterministic vs. Smoothing

table III compares AEGIS’s deterministic per-node certificates (section III) against randomized smoothing [3] ($\sigma = 0.01$, 200 Monte Carlo samples, Clopper-Pearson confidence $1 - \alpha = 0.999$).

TABLE III: Deterministic (AEGIS) vs. smoothing certificates (10 seeds). Equal coverage; AEGIS radii 2–5 $\times$ larger.

Dataset	Coverage		Median radius	
	Det	Smooth	Det	Smooth
Cora	81 $\pm$ 9%	81 $\pm$ 9%	.056 $\pm$.02	.015 $\pm$.00
Citeseer	83 $\pm$ 4%	83 $\pm$ 4%	.078 $\pm$.06	.015 $\pm$.00
WikiCS	70 $\pm$ 3%	70 $\pm$ 3%	.030 $\pm$.01	.015 $\pm$.00

TABLE IV: Scalability of AEGIS analysis on Cora ego-subgraphs.

N	$ E $	D	J_z	J_A	Total
20	22	1,280	0.2s	0.2s	0.5s
50	61	3,200	0.5s	0.7s	1.3s
100	158	6,400	1.1s	1.6s	2.8s
200	389	12,800	2.3s	5.1s	8.1s

At matched coverage, AEGIS certificates provide radii 2.0–5.2 $\times$ larger than smoothing. The smoothing radius is constant (≈ 0.015) because it depends on σ and sample count, not the model’s structural sensitivity. AEGIS certificates are deterministic (always valid), while smoothing holds with probability $1 - \alpha$.

D. Phase Transition Validation

To validate section III, we scale the adjacency matrix to achieve target spectral radii $\rho \in \{0.3, 0.5, 0.7, 0.85, 0.9, 0.95, 0.99\}$ and measure the actual fixed-point shift under a random perturbation of $\varepsilon = 0.01$. The shift amplifies by 83 $\times$ as ρ increases from 0.3 to 0.99, confirming the three-regime behavior: bounded shift in the subcritical regime, rapid growth near $\rho = 1$, and divergence in the supercritical regime.

E. Scalability

table IV reports wall-clock time for the full AEGIS pipeline on Cora ego-subgraphs of varying size (single NVIDIA RTX 4090 GPU).

The practical GPU limit is $N \approx 200$ ($D = 12,800$). Empirical scaling is 16 $\times$ for 10 $\times$ growth in D , dominated by the $O(D \cdot N^2)$ numerical Jacobian J_A rather than the $O(D^3)$ linear solve. This covers all IEEE test cases (14–118 buses) and practical ego-network sizes for node-level analysis.

F. Case Study: Power Flow N-1 Contingency

As a cross-domain application, we train ContractiveGCN-PF on IEEE benchmark grids and compare the AEGIS vulnerability spectrum against brute-force N-1 contingency analysis (remove each line, re-solve, measure deviation).

Tightness of ≈ 1.00 confirms that exact first-order prediction transfers to the power flow domain. N-1 ranking correlation is strongest on larger grids (case57: $\tau = +0.67$, case118: $\tau = +0.62$), where edges exhibit more diverse criticalities for IFT to differentiate. This demonstrates that the adversarial vulnerability spectrum and power grid contingency

TABLE V: AEGIS on IEEE power flow benchmarks (10 seeds). Tightness ≈ 1.00 ; N-1 ranking strongest on larger grids.

Case	N	$ E $	Tight.	N-1 τ	Top-5
case14	14	20	$1.00 \pm .01$	$+ .42 \pm .19$	$70 \pm 10\%$
case30	30	41	$1.00 \pm .01$	$+ .37 \pm .17$	$66 \pm 13\%$
case57	57	78	$1.02 \pm .01$	$+ .67 \pm .09$	$52 \pm 24\%$
case118	118	179	$1.01 \pm .01$	$+ .62 \pm .11$	$62 \pm 14\%$

ranking are instances of the same mathematical object—the structural sensitivity matrix S .

VI. RELATED WORK

Adversarial attacks on GNNs. Structural attacks modify graph topology to degrade GNN performance. Nettack [1] targets individual nodes via greedy edge flips; Mettack [2] uses meta-gradients through a surrogate for global poisoning. These methods operate on explicit (finite-depth) GNNs and require surrogate transfer when applied to DEQ models. AEGIS directly analyzes the equilibrium via IFT, yielding $3\text{--}10\times$ stronger attacks without surrogates.

Certified robustness for GNNs. Randomized smoothing [3] provides probabilistic certificates by classifying under random perturbation of the graph structure. Interval bound propagation has been applied to GNNs [10] but requires layer-by-layer propagation, incompatible with infinite-depth DEQ models. AEGIS provides deterministic certificates via the IFT with $2\text{--}5\times$ larger radii than smoothing at equal coverage.

DEQ robustness and Lipschitz analysis. Lipschitz bounds for DEQs have been studied for input perturbation [8], [9], [11]. Pabbaraju et al. [12] estimate Lipschitz constants of monotone DEQs. These works bound $\|\partial z^*/\partial x\|$ but not $\|\partial z^*/\partial A\|$ —the structural sensitivity central to graph adversarial robustness. AEGIS fills this gap with the structural sensitivity matrix S and the phase transition characterization.

Implicit GNNs. IGNN [5] enforces contractivity via spectral normalization. EIGNN [13] improves computational efficiency. Neither analyzes adversarial structural vulnerability at the equilibrium. The concurrent work of Liu et al. [14] studies stability of GNNs under graph perturbation but without the phase transition framework or certified per-node radii.

VII. CONCLUSION

We presented AEGIS, a framework that exploits the equilibrium structure of DEQ-GNNs for exact adversarial vulnerability prediction. Our phase transition theorem characterizes three regimes of structural vulnerability around a sharp critical budget $\varepsilon_{\text{crit}}$. The constrained sensitivity analysis achieves tightness 1.00 ± 0.01 across 9 datasets—the first-order prediction is essentially exact for realistic graph perturbations. AEGIS attacks dominate Mettack (30/30 wins), its deterministic certificates are $2\text{--}5\times$ larger than smoothing, and the vulnerability spectrum bridges adversarial ML with power systems contingency analysis through a single mathematical object.

Limitations and future work. Dense Jacobian computation limits subgraph size to $N \approx 200$. The first-order approximation is exact only for small perturbations; finite perturbations incur higher-order gaps (tightness degrades for $\varepsilon > 0.1$). Extending to sparse Neumann-series solvers would enable larger-scale analysis. Applying AEGIS to other implicit architectures (monotone operator networks, neural ODEs on graphs) and discrete perturbation models (edge insertion/deletion) are natural next steps.

REFERENCES

- [1] D. Zügner, A. Akbarnejad, and S. Günnemann, “Adversarial attacks on neural networks for graph data,” in *Proceedings of the 24th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, 2018, pp. 2847–2856.
- [2] D. Zügner and S. Günnemann, “Adversarial attacks on graph neural networks via meta learning,” in *International Conference on Learning Representations*, 2019.
- [3] A. Bojchevski, J. Klicpera, and S. Günnemann, “Efficient robustness certificates for discrete data: Sparsity-aware randomized smoothing for graphs, images and more,” in *International Conference on Machine Learning*, 2020, pp. 1003–1013.
- [4] S. Bai, J. Z. Kolter, and V. Koltun, “Deep equilibrium models,” in *Advances in Neural Information Processing Systems*, vol. 32, 2019.
- [5] F. Gu, H. Chang, W. Zhu, S. Sojoudi, and L. El Ghaoui, “Implicit graph neural networks,” in *Advances in Neural Information Processing Systems*, vol. 33, 2020, pp. 11 984–11 995.
- [6] J. Lorraine, P. Vicol, and D. Duvenaud, “Optimizing millions of hyperparameters by implicit differentiation,” in *International Conference on Artificial Intelligence and Statistics*, 2020, pp. 1540–1552.
- [7] S. Gould, R. Hartley, and D. Campbell, “Deep declarative networks,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 44, no. 8, pp. 3988–4004, 2021.
- [8] L. El Ghaoui, F. Gu, B. Travacca, A. Askari, and A. Tsai, “Implicit deep learning,” *SIAM Journal on Mathematics of Data Science*, vol. 3, no. 3, pp. 930–958, 2021.
- [9] M. Revay, R. Wang, and I. R. Manchester, “Lipschitz bounded equilibrium networks,” in *arXiv preprint arXiv:2010.01732*, 2020.
- [10] D. Zügner and S. Günnemann, “Certifiable robustness and robust training for graph convolutional networks,” in *Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, 2019, pp. 246–256.
- [11] E. Winston and J. Z. Kolter, “Monotone operator equilibrium networks,” in *Advances in Neural Information Processing Systems*, vol. 33, 2020, pp. 10 718–10 728.
- [12] C. Pabbaraju, E. Winston, and J. Z. Kolter, “Estimating lipschitz constants of monotone deep equilibrium models,” *arXiv preprint arXiv:2111.08931*, 2021.
- [13] Y. Yang, J. Tang, and C. Luo, “Efficient computation of higher-order subgraph attribution via message passing,” in *International Conference on Machine Learning*, 2022.
- [14] S. Liu, R. Ying, H. Dong, L. Li, T. Xu, Y. Liang, J. Leskovec, and H. Peng, “Stability and generalization of graph neural networks,” in *arXiv preprint*, 2021.